

Figure and table legends

Provisional companion file. The same numbered figures, tables, and captions remain embedded in the blinded manuscript.

Figure 1. Composable assurance and the evidence ladder used throughout the paper.

Figure 2. Technology-neutral architecture and the sovereign boundary.

Figure 3. Cross-border sequence with explicit PREPARE and terminal decision certificates.

Figure 4. Institutional separation, learned privacy attacks and registered public-statistics releases.

Figure 5. Learned pairwise-attack ROC-AUC under unseen-user and independently generated evaluation regimes.

Figure 6. Utility and clipping effects for dense and sparse registered histograms.

Figure 7. Leakage-controlled AML comparison and prospective effectiveness outcomes.

Figure 8. Mean PR-AUC difference by independent generator family.

Figure 9. Sanctions screening from official-list ingestion to analyst and ownership decisions.

Figure 10. Calibrated sanctions policy frontier: recall and false-positive trade-off.

Figure 11. Prevalence-adjusted analyst workload across calibrated policies.

Figure 12. From jurisdictional data to a partially identified and adaptively governed policy set.

Figure 13. Illustrative holding-limit ceilings across exposure tolerances, eligible populations and adoption shares.

Figure 14. Four-stage national-scale qualification and the claim boundary.

Figure 15. Trace-driven digital-twin p99 latency across workload and failure scenarios.

Figure 16. Analytical lower-bound units by stage and qualification workload.

Table 1. Principal novelty boundaries.

Table 2. Core design requirements.

Table 3. Cross-border protocol objects.

Table 4. Cross-border properties and proof boundary.

Table 5. Privacy attacker model.

Table 6. Learned network-observer linkage results.

Table 7. Ranking attack under independently generated traffic.

Table 8. Differential-privacy utility by annual budget.

Table 9. AML/CFT predictive-evaluation design.

Table 10. Family-specific AML effects.

Table 11. Predictive workload at fixed alert budgets.

Table 12. Sanctions evaluation design.

Table 13. Calibrated screening policies.

Table 14. Prevalence-adjusted workload illustration.

Table 15. Recall by evidence and perturbation class.

Table 16. Illustrative candidate limits under one exposure definition.

Table 17. Mechanical duration imposed by system-wide flow limits (days).

Table 18. Digital-twin scenario outcomes.

Table 19. Necessary parallel-unit lower bounds (2x headroom, 60% planned utilization).

Table 20. Mandatory physical and pilot acceptance evidence.

Table 21. Security controls and evidence.

Table 22. Claims-to-evidence register.